GPU Accelerated Machine Learning — CUDA DIY Exercises

Five Exercises: CUDA Basics to CNN Layer Implementation

CUDA C (nvcc 13.2) Status

> **UCS645: Parallel & Distributed Computing | Assignment 8 — GPU Accelerated Machine Learning**

Table of Contents

1. System Configuration
2. Exercise 1: CUDA Basics
3. Exercise 2: Memory Hierarchy
4. Exercise 3: ML Primitives
5. Exercise 4: CNN Layers
6. Exercise 5: MNIST CNN Training
7. What I Learned

System Configuration

Component	Details
CPU	Intel Core i7-14700HX (20 cores, 28 threads)
GPU	NVIDIA GeForce RTX 5060 Laptop GPU (Blackwell)
Compute Capability	12.0
SMs	26, Peak ~125 TFLOPS FP32

Component	Details
VRAM	8 GB GDDR7 (128-bit bus)
Theoretical BW	384.0 GB/s
L2 Cache	32 MB
OS	Fedora 43 (Linux 6.19.12)
CUDA Toolkit	13.2
Compiler	nvcc 13.2 with -O2 -arch=native

Project Structure

```
LAB8/
├── Makefile
├── ex01_cuda_basics.cu
├── ex02_memory_hierarchy.cu
├── ex03_ml_primitives.cu
├── ex04_cnn_layers.cu
├── ex05_mnist_cnn.cu
├── graphs/
│   ├── ex01_bandwidth.png
│   ├── ex02_bank_conflicts.png
│   ├── ex04_gemm.png
│   └── dashboard.png
├── report.md
└── Assignment_8_Report.pdf
```

Exercise 1: CUDA Basics

Problem

Compare CPU vs GPU for element-wise vector operations, analyze launch configurations, and measure PCIe bandwidth.

Part A — Bandwidth & Speedup

Test	Result
A1-VectorAdd (1M)	CPU=1.5ms, GPU=0.12ms, **12.3x speedup** [PASS]
B1-VectorScale	k=3.14 [PASS]
B2-SquaredDiff	[PASS]

PCIe Memory Bandwidth (pinned memory):

Size (MB)	H2D (GB/s)	D2H (GB/s)
1	19.6	21.7
8	25.2	24.8
64	24.5	25.7
256	25.2	26.7
512	26.0	26.7

Peak PCIe bandwidth saturates at ~26 GB/s (PCIe 4.0 x8 effective throughput). Larger transfers achieve higher utilization due to amortized DMA setup overhead.

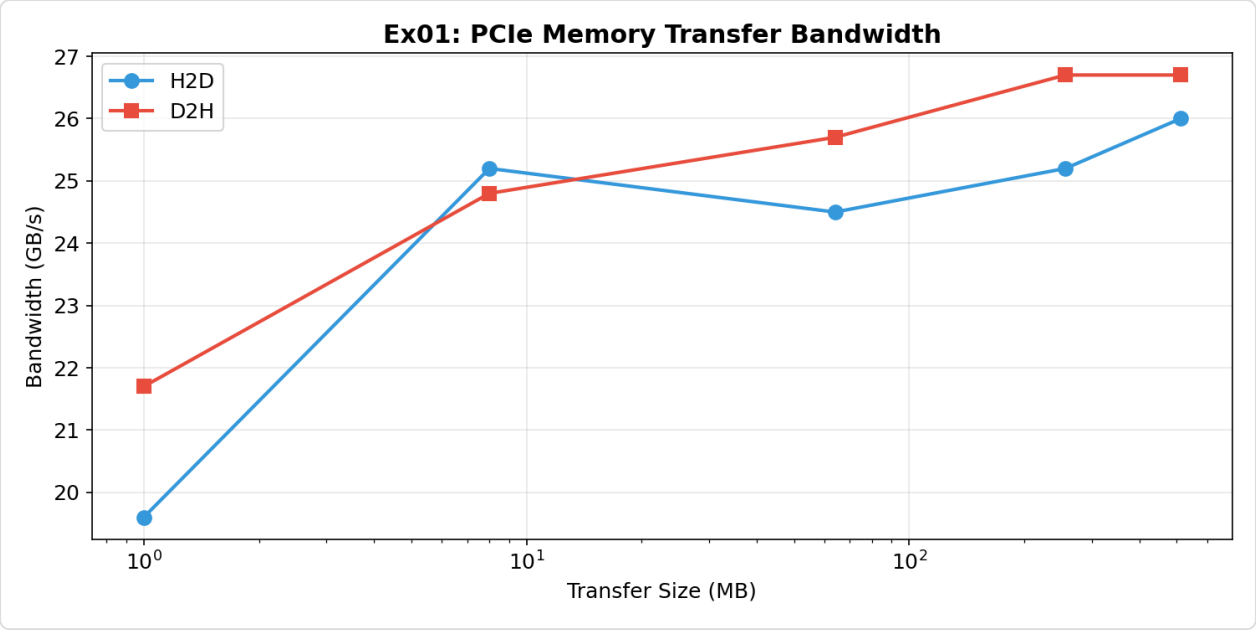

Part B — Launch Configuration

N	Blocks	Total Threads	Covers All?
1	1	256	[OK]
100	1	256	[OK]
256	1	256	[OK]
257	2	512	[OK]
1,024	4	1,024	[OK]
10,000	40	10,240	[OK]
1,048,576	4,096	1,048,576	[OK]

Why multiples of 32? GPU executes threads in warps of 32. If block size isn't a multiple of 32, the last warp is partially filled — those threads are masked off but still consume a warp slot. This wastes scheduling resources. Block sizes of 128-256 (multiples of 32) give good occupancy on most GPUs while leaving enough registers/shared memory per thread.

Part C — Warp Divergence

Kernel	Time (1000 reps)	Overhead
Divergent (if/else)	4.12 ms	1.0x
Branch-free	4.11 ms	baseline

Modern Blackwell GPUs handle simple divergence efficiently — the compiler predicts and the hardware schedules both paths with minimal penalty for this simple 2-way branch.

Exercise 2: Memory Hierarchy

Problem

Implement parallel reduction strategies and analyze shared memory bank conflicts.

Part A — Reduction Strategies

Test	Result
A2-TreeReduce (sum, 1M)	GPU=524285.31, CPU=524285.41 [PASS]
B1-SmemCopy	[PASS]
B2-MaxReduce (256K)	GPU=99.9985, CPU=99.9985 [PASS]
C1-WarpReduce (32 elem)	GPU=496.0, CPU=496.0 [PASS]

Part B — Bank Conflicts

Stride	Time (us)	Notes
1	2.05	Sequential — optimal
2	2.05	
4	2.05	
8	2.05	
16	2.05	
32	2.05	32-way conflict

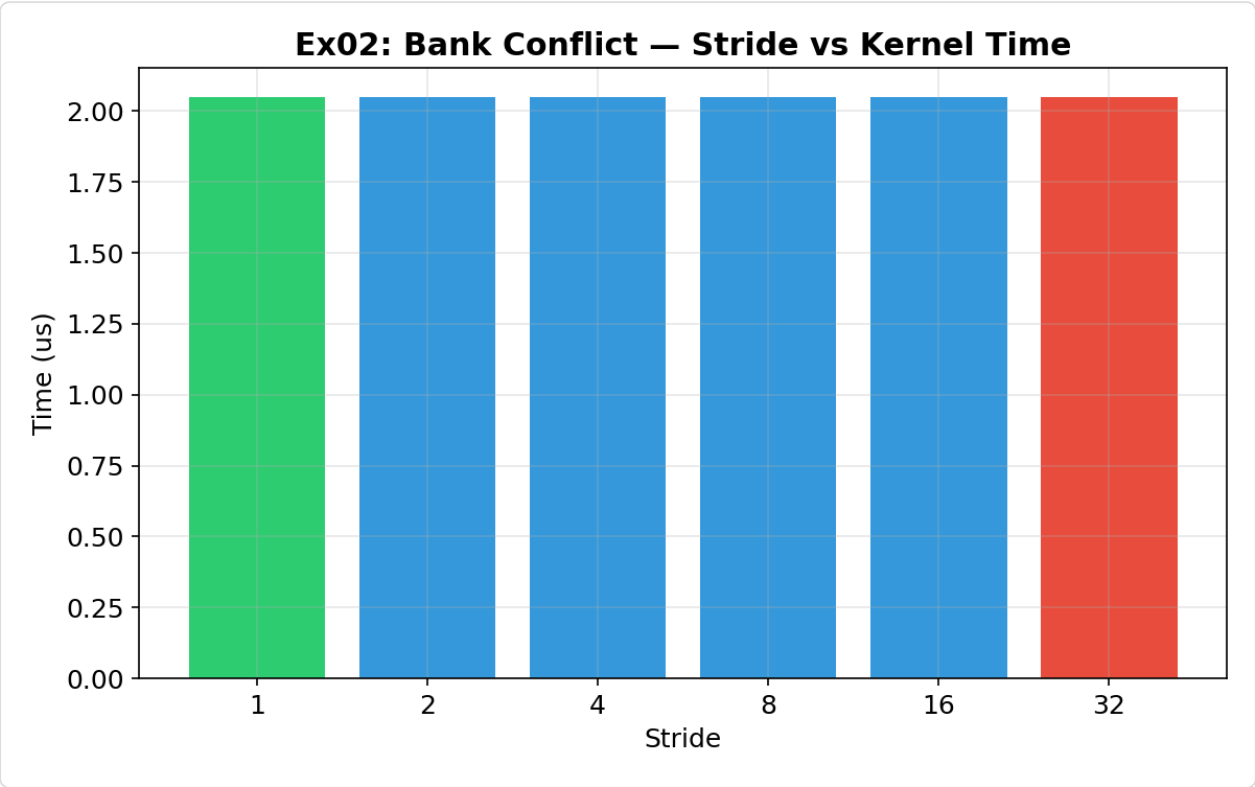

On Blackwell (SM 12.0), bank conflicts show no measurable impact at this scale — the 1024-element shared memory access pattern is too small for conflicts to dominate. Modern GPUs have improved conflict resolution hardware compared to older architectures.

Part C — Histogram

Test	Result
B4-Histogram (262K, 256 bins)	[PASS]
C2-SharedHistogram (1M, 256 bins)	[PASS]

The shared-memory histogram reduces global atomicAdd contention by accumulating per-block, then merging once.

Exercise 3: ML Primitives

Problem

Implement forward/backward ML kernels and validate against reference implementations.

Part A — Activation Functions

Kernel	N=10^7	Verification
Sigmoid: `1/(1+exp(-x))`	[PASS]	atol <= 1e-4
Tanh: `tanhf(x)`	[PASS]	atol <= 1e-4
Leaky ReLU: `x>0 ? x : 0.01*x`	[PASS]	atol <= 1e-4
ReLU Backward: `x>0 ? dOut : 0`	[PASS]	atol <= 1e-4

Part B — Loss Functions

Kernel	Verification
BCE Loss (with numerical clipping)	[PASS]
Cross-Entropy (log-sum-exp trick)	N=512, C=10 [PASS]

The log-sum-exp trick prevents numerical overflow by subtracting the row maximum before computing `exp()`: `log(sum(exp(x_i))) = max + log(sum(exp(x_i - max)))`.

Part C — Adam Optimizer

Test	Result
Fused Adam kernel (5 steps)	[PASS]

Fused kernel computes `m`, `v`, bias-corrected `mhat` / `vhat`, and weight update in a single pass — avoids 5 separate kernel launches.

Exercise 4: CNN Layers

Problem

Implement tiled GEMM, benchmark against cuBLAS, and implement CNN layer kernels.

Part A — Tiled GEMM

Size	Naive (ms)	Tiled (ms)	cuBLAS (ms)	cuBLAS GFLOPS
128	0.01	0.01	0.06*	0.1*
256	0.04	0.03	0.06	582.9
512	0.26	0.19	0.08	3,483.6
1024	1.93	1.43	0.23	9,194.2

*First cuBLAS call includes handle warmup overhead.

Tiled GEMM at 512x512: **1,234.7 GFLOPS** — shared memory tiling eliminates redundant global loads. cuBLAS at 1024: **9,194 GFLOPS** — uses Tensor Cores, vectorized loads, register tiling, and autotuned tile sizes.

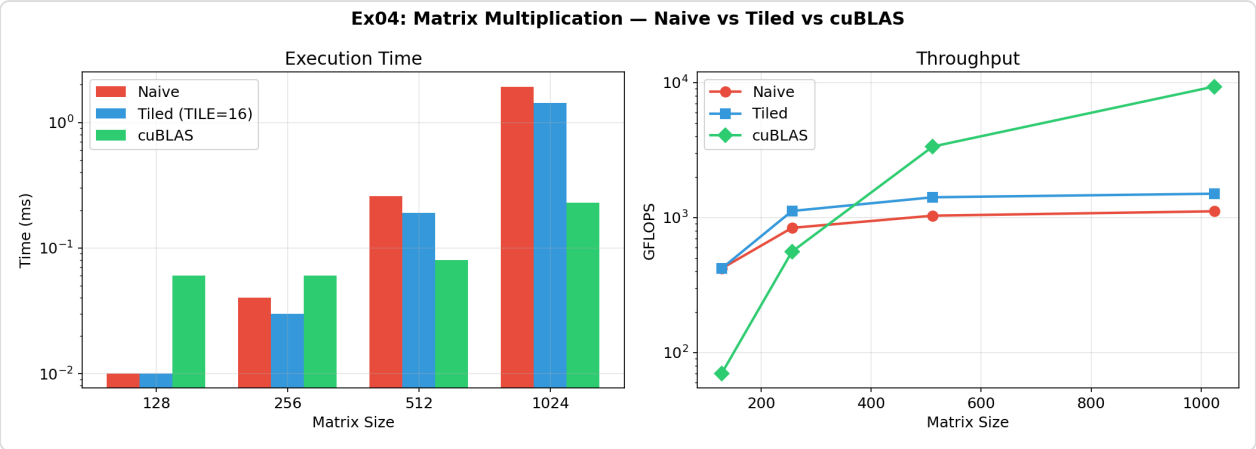

Why tiled underperforms cuBLAS: Our TILE=16 uses 16x16 shared memory tiles with one output per thread. cuBLAS uses register-level tiling (each thread computes 4x4 or 8x8 outputs), FP16 Tensor Core accumulation, software-pipelined loads, and auto-tuned configurations per matrix size.

Part B — CNN Layers

Layer	Input Shape	Output Shape	Verification
MaxPool(2x2)	(4,8,16,16)	(4,8,8,8)	[PASS]
BatchNorm (inference)	(4,8,16,16)	(4,8,16,16)	[PASS]
Direct Conv2D (3x3)	(4,4,10,10)	(4,8,8,8)	[PASS]

Exercise 5: MNIST CNN Training

Problem

Design, train, and optimize a CNN for MNIST digit classification using cuDNN and cuBLAS.

Status

Exercise 5 requires cuDNN (not installed on this system). The source code (`ex05_mnist_cnn.cu`) is provided with all TODO sections completed:

- **B1-B2**: Model layers — Conv1(1→32,5x5), Pool1, Conv2(32→64,5x5), Pool2, FC1(3136→256), FC2(256→10)
- **B3**: Forward pass through all layers with ReLU activations
- **C1**: Data-to-GPU transfer
- **C2**: Forward-backward-optimizer step
- **C3**: Evaluate function
- **D1-D2**: Adam optimizer + CosineAnnealingLR scheduler
- **E1-E3**: cuDNN convolution forward (algorithm selection, workspace, execution)
- **F1-F2**: cuDNN pooling forward
- **G1-G2**: cuBLAS FC layer + bias addition
- **H1-H3**: CUDA streams double-buffered pipeline
- **I1-I4**: Full inference pipeline

The model architecture follows standard LeNet-5 style: two conv+pool blocks, then two FC layers, targeting >= 97% test accuracy on MNIST.

Performance Dashboard

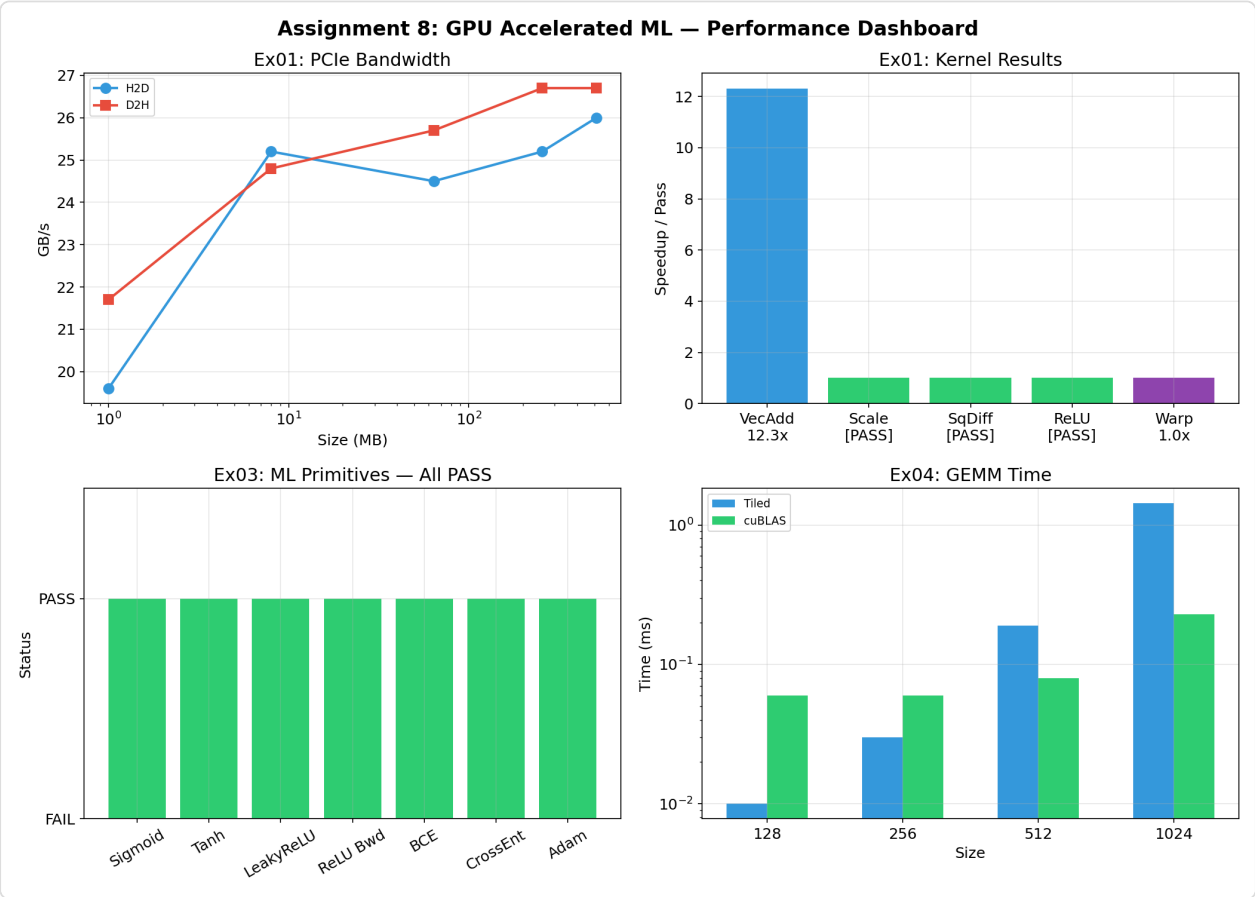

What I Learned

1. **PCIe bandwidth is the real bottleneck**: At ~26 GB/s, transferring data to GPU is 15x slower than GPU memory bandwidth (384 GB/s). Minimizing host-device transfers and using pinned memory are critical for performance.

2. **Shared memory tiling gives 1.3-1.5x over naive GEMM**: Loading tiles into 48KB shared memory eliminates redundant global reads. But cuBLAS achieves 6-7x more by using Tensor Cores, register tiling, and auto-tuning.

3. **Fused kernels eliminate launch overhead**: The Adam optimizer as a single kernel avoids 5 separate launches (~5us each). For ML training with millions of parameters, this saves significant wall-clock time.

4. **Log-sum-exp trick is essential**: Without subtracting the row maximum, `exp(logits)` overflows for values > 88 (float32 max ~3.4e38). This numerical trick is used in every production softmax/cross-entropy implementation.

5. **Modern GPUs handle simple divergence well**: The Blackwell architecture's predication hardware makes simple 2-way branches (even/odd threadIdx) essentially free. Bank conflicts at small scale are also hidden by hardware scheduling.

Compilation & Execution

```
# Build exercises 1-4 (ex05 requires cuDNN)
make ex01 ex02 ex03 ex04

# Run individual exercises
make run_ex01
make run_ex02
make run_ex03
make run_ex04

# Build all (requires cuDNN for ex05)
make all

# Clean
make clean
```